Day 1: Introduction to Data Science

Kenneth Benoit and Slava Mikhaylov

Introduction to Data Science and Big Data Analytics

17 August 2015

Day 1 Outline

- Logistics and introductions
- Defining Data Science
- The R statistical package
- Data structures
- Further course preview
- Preview of lab sessions

Pitch of this course

- ▶ Whom this class is for
  - ▶ (typically) graduate-level applied researchers with some prior experience in quantitative methods
  - ▶ those requiring the fundamentals of data science
  - ▶ people working with typically large datasets and databases
  - ▶ those seeking a practical introduction to data analysis using R

- ▶ Prerequisites
  - ▶ at least one prior course in quantitative methods or statistics
  - ▶ familiarity, or willingness to learn, R
  - ▶ ability to use a text editor and spreadsheet is helpful
  - ▶ (optional) ability to process data files in a programming language such as Python

Learning objectives

- an overview of data science and the challenges of working with big data

- learning the R statistical package and practical methods for manipulating and analyzing data

- formulating research questions using data analytics

- how to acquire, process, store, and use data, both structured and unstructured

- the basics of statistical inference including probability and probability distributions, modelling, experimental design

Learning objectives (cont.)

- an overview of classification methods and methods for assessing model fit

- supervised learning approaches, including linear and logistic regression, decision trees, and naive Bayes

- unsupervised learning approaches, including clustering, association rules, and principal components analysis

- quantitative methods of text analysis, including mining social media and other online resources

Who we are

- In the lectures:
  - Ken Benoit, London School of Economics
    kbenoit@lse.ac.uk
  - Slava Mikhaylov, University College London
    s.mikhaylov@ucl.ac.uk

- In the labs:
  - Paul Nulty, p.nulty@lse.ac.uk

- Course homepage: `http://github.com/kbenoit/ME114`

Class schedule: Typical day

10:00–11:15 Lecture part 1 (NAB Thai Theatre)

11:25–12:45 Lecture part 2 (NAB Thai Theatre)

14:00–15:30 Lab Group 1 (TW2.4.02)

15:30–17:00 Lab Group 2 (TW2.4.02)

Essential course resources

- On-line course resources
  GitHub repository `http://github.com/kbenoit/ME114`
  Moodle page `https://shortcourses.lse.ac.uk/course/`
              `view.php?id=158`
- Using GitHub and forking the course repo
- RStudio and GitHub integration
- How to complete (and submit) assignments
  - For a brief on-line introduction to RMarkdown, which we will use for completing the exercises for the course, see `https://goo.gl/ZqOwUe`

Course reading materials

- The R statistical package resources on https://cran.r-project.org/doc/manuals
- Key readings
  - James et al. (2013) *An Introduction to Statistical Learning: With applications in R*. Springer.
  - Zumel, N. and Mount, J. (2014). *Practical Data Science with R*. Manning Publications.
- Additional readings – recommended but we understand that this course will be intensive

What is "data science"?

- extraction or generation of knowledge from data
- extends "data mining", using computational and algorithmic methods
- combines applied statistical methods with advances in computer science, especially machine learning
- may involve "unstructured" data, especially text, but also video and images
- closely tied to computational methods

Why focus on data types and structures?

- "data" mining and "data" science imply that we know how to work with data
- data structures are not neutral - they shape how we record, see, and have the ability to analyze information
- much of the actual work in data mining and data analysis is done at the data "mungeing" stage

Basic (atomic) data types in R

numeric 8-byte numeric representations

integer non-floating point numbers

character text

logical TRUE or FALSE

Recursive types also exist, such as lists and vectors; there are also special classifications for NA

Basic data types in R: integer

```r
x <- 10
typeof(x)

## [1] "double"

is.integer(x)

## [1] FALSE

x <- 7L          # force integer type
typeof(x)

## [1] "integer"

object.size(x)

## 48 bytes

as.integer(3.14)

## [1] 3
```

Basic data types in R: character

```r
typeof("test string")

## [1] "character"

object.size("a")

## 96 bytes

s <- ""              # Unicode
cat(s)

##

as.character("3.14")   # coerce numerics to character

## [1] "3.14"
```

Basic data types in R: numeric

```r
x <- 10.5        # assign a numeric value
x                # print the value of x

## [1] 10.5

typeof(x)         # print the class name of x

## [1] "double"

object.size(x)   # show storage size in bytes

## 48 bytes
```

is.*() and as.*()

```r
is.numeric(x)     # is the object of numeric type?

## [1] TRUE

is.numeric(7.1)

## [1] TRUE

is.numeric("7.1")

## [1] FALSE

is.numeric(as.numeric("7.1"))

## [1] TRUE
```

Basic data types in R: logical

A logical value is 'TRUE' or 'FALSE', often created via comparison between variables.

```
1 < 2                # is 1 less 2

## [1] TRUE

x <- c(1, 2, 3)
y <- c(4, 3, 2)
x > y                # vectorized comparison

## [1] FALSE FALSE  TRUE

typeof(x > y)

## [1] "logical"
```

Difference between 'mode' and 'class'

- ▶ 'atomic' modes are numeric, complex, character and logical
- ▶ recursive objects have modes such as 'list' or 'function' or a few others
- ▶ an object has one and only one mode
- ▶ 'class' is a property assigned to an object that determines how generic functions operate with it - not a mutually exclusive classification
- ▶ an object has no specific class assigned to it, such as a simple numeric vector, it's class is usually the same as its mode, by convention
- ▶ an object's mode can be changed through coercion, without necessarily changing the class

Numerical precision issues

- floating point numbers are approximations of numbers
  - precision: anything more than 16 base-10 digits must be approximated
  - fractions: approximated if not $\frac{x}{2^k}$
  - anything over stated precision is truncated: $3.57e21 + 1 = 3.57e21$

```
1 - 4/5 - 1/5   # not zero!
## [1] -5.551115e-17
```

Machine limits

```
.Machine$integer.max

## [1] 2147483647

.Machine[c("double.xmin", "double.xmax", "double.digits")]

## $double.xmin
## [1] 2.225074e-308
##
## $double.xmax
## [1] 1.797693e+308
##
## $double.digits
## [1] 53
```

Alternatives (Stata)

- single and double precision:
  http://blog.stata.com/2012/04/02/
  the-penultimate-guide-to-precision/
- R has only double precision

Common input formats

- csv
- Excel
- "fixed formats"
- relational databases
- embedded tags: Extensible Markup Language (XML)
- key-value pair schemes (JSON)
    - examples of JSON and XML: `http://json.org/example`

Special issue: text encoding

- a "character set" is a list of character with associated numerical representations
- ASCII: the original character set, uses just 7 bits (2^7) – see `http://ergoemacs.org/emacs/unicode_basics.html`
- ASCII was later extended, e.g. ISO-8859 `http://www.ic.unicamp.br/~stolfi/EXPORT/www/ISO-8859-1-Encoding.html`, using 8 bits (2^8)
- but this became a jungle, with no standards: `http://en.wikipedia.org/wiki/Character_encoding`

Solution: Unicode

- Unicode was developed to provide a unique number (a "code point") to every known character – even some that are "unknown"
- problem: there are more far code points than fit into 8-bit encodings. Hence there are multiple ways to *encode* the Unicode code points
- *variable-byte* encodings use multiple bytes as needed. Advantage is efficiency, since most ASCII and simple extended character sets can use just one byte, and these were set in the Unicode standard to their ASCII and ISO-8859 equivalents
- two most common are UTF-8 and UTF-16, using 8 and 16 bits respectively

Warnings with text encodings

- Input texts can be very different
- Many text production software (e.g. MS Office-based products) still tend to use proprietary formats, such as Windows-1252
- Windows tends to use UTF-16, while Mac and other Unix-based platforms use UTF-8
- Your eyes can be deceiving: a client may display gibberish but the encoding might still be as intended
- No easy method of detecting encodings (except in HTML meta-data)

What is a "Dataset"?

- A dataset is a "rectangular" formatted table of data in which all the values of the same variable must be in a single column
- Many of the datasets we use have been artificially reshaped in order to fulfill this criterion of rectangularity

Revisting basic data concepts

- The difference between tables and *datasets*
- This is a (partial) dataset:

```
district      incumbf wonseatf
1        Carlow Kilkenny Challenger      Lost
2        Carlow Kilkenny Challenger      Lost
5        Carlow Kilkenny  Incumbent       Won
100 Donegal South West Challenger      Lost
459             Wicklow  Incumbent       Won
464             Wicklow Challenger      Lost
```

- This is a table:

```
           Lost Won
Challenger  266  60
Incumbent    32 106
```

- The key with a dataset is that all the values of the same variable must be in a single column

Example: Comparative Manifesto Project dataset

Note: Available from `https://manifestoproject.wzb.eu/`

```r
# load in a subset of the Manifesto Project dataset, with counts
load(url("http://kenbenoit.net/files/cmpdata.Rdata"))
# View(cmpdata)
```

Example: Comparative Manifesto Project dataset

This is "wide" format:

	country	countryname	oecdmember	eumember	edate	date	party	partyname
203	42	Austria	10	10	2008-09-28	200809	42320	SPOE Social Democratic Party
204	42	Austria	10	10	2008-09-28	200809	42110	Green Party
205	42	Austria	10	10	2008-09-28	200809	42520	OVP: People's Party
206	42	Austria	10	10	2008-09-28	200809	42420	FPO: Freedom Party
207	42	Austria	10	10	2008-09-28	200809	42710	BZO Alliance for the Future of Austria
208	42	Austria	10	10	2008-09-28	200809	42220	KP<d6> Communist Party of Austria
314	21	Belgium	10	10	1991-11-24	199111	21521	CVP Christian People's Party
315	21	Belgium	10	10	1991-11-24	199111	21111	ECOLO Francophone Ecologists
316	21	Belgium	10	10	1991-11-24	199111	21321	SP Flemish Socialist Party
317	21	Belgium	10	10	1991-11-24	199111	21522	PSC Christian Social Party
318	21	Belgium	10	10	1991-11-24	199111	21421	PVV Party of Liberty and Progress
319	21	Belgium	10	10	1991-11-24	199111	21913	VU People's Union
320	21	Belgium	10	10	1991-11-24	199111	21912	FDF Francophone Democratic Front

Long v. wide formats

- reshape
  - the "old" R way to do this, using 'base::reshape()'
  - problem: confusing and difficult to use
- reshape2
  - from Hadley Wickham's `reshape2` package
  - data is first 'melt'ed into long format
  - then 'cast' into desired format

Example: wide to long using `reshape2`

```r
require(reshape2, quietly = TRUE)
# this will select only the "per" variables for measurement
cmpdataLong <- melt(cmpdata,
id.vars = c("countryname", "party", "date"),
measure.vars = names(cmpdata)[21:76],
variable.name = "category",
value.name = "catcount")
# why do this?
cmpdataLong$category <- as.character(cmpdataLong$category)
# View(cmpdataLong)
```

Example: wide to long using `reshape2`

```
require(reshape2, quietly = TRUE)
# now we can get summary statistics across countries, e.g. for economic
with(subset(cmpdataLong, grepl("^per7", category)),
     table(countryname, category))
```

```
##                  category
## countryname       per701 per702 per703 per704 per705 per706
##     Austria          34     34     34     34     34     34
##     Belgium          63     63     63     63     63     63
##     Cyprus           10     10     10     10     10     10
##     Denmark          60     60     60     60     60     60
##     Finland          47     47     47     47     47     47
##     France           23     23     23     23     23     23
##     Germany          30     30     30     30     30     30
##     Great Britain    20     20     20     20     20     20
##     Greece           17     17     17     17     17     17
##     Iceland          31     31     31     31     31     31
##     Ireland          31     31     31     31     31     31
##     Israel           32     32     32     32     32     32
##     Italy            41     41     41     41     41     41
##     Luxembourg       21     21     21     21     21     21
##     Malta             4      4      4      4      4      4
##     Netherlands      48     48     48     48     48     48
##     Norway           28     28     28     28     28     28
##     Portugal         38     38     38     38     38     38
```

A better way

```
require(dplyr, warn.conflicts = FALSE, quietly = TRUE)
cmpdataLong2 <- melt(select(cmpdata, countryname, party, date, per101:per706),
                     id.vars = c("countryname", "party", "date"),
                     # NOT NEEDED measure.vars = names(cmpdata)[21:76],
                     variable.name = "category",
                     value.name = "catcount")
cmpdataLong2$category <- as.character(cmpdataLong2$category)
identical(cmpdataLong, cmpdataLong2)

## [1] TRUE
```

A better way

```r
with(filter(cmpdataLong2, grepl("^per7", category)),
     table(countryname, category))
```

```
##                 category
## countryname      per701 per702 per703 per704 per705 per706
##    Austria           34     34     34     34     34     34
##    Belgium           63     63     63     63     63     63
##    Cyprus            10     10     10     10     10     10
##    Denmark           60     60     60     60     60     60
##    Finland           47     47     47     47     47     47
##    France            23     23     23     23     23     23
##    Germany           30     30     30     30     30     30
##    Great Britain     20     20     20     20     20     20
##    Greece            17     17     17     17     17     17
##    Iceland           31     31     31     31     31     31
##    Ireland           31     31     31     31     31     31
##    Israel            32     32     32     32     32     32
##    Italy             41     41     41     41     41     41
##    Luxembourg        21     21     21     21     21     21
##    Malta              4      4      4      4      4      4
##    Netherlands       48     48     48     48     48     48
##    Norway            28     28     28     28     28     28
##    Portugal          38     38     38     38     38     38
##    Spain             57     57     57     57     57     57
##    Sweden            44     44     44     44     44     44
```

Grouping operations: number of parties per election

```
# group by country-election
by_country <- group_by(cmpdataLong, countryname, date)
nparties <- summarise(by_country, npart = n())
head(nparties)

## Source: local data frame [6 x 3]
## Groups: countryname
##
##   countryname   date npart
## 1     Austria 199010   224
## 2     Austria 199410   280
## 3     Austria 199512   280
## 4     Austria 199910   224
## 5     Austria 200211   280
## 6     Austria 200610   280

# is that correct?
```

Grouping operations: number of parties per election corrected

```r
# group by country-election
by_country_unique <- distinct(cmpdataLong, countryname, date, party)
by_country_n <- group_by(by_country_unique, countryname, date)
nparties <- summarise(by_country_n, npart = n())
head(nparties, 10)

## Source: local data frame [10 x 3]
## Groups: countryname
##
##    countryname   date npart
## 1      Austria 199010     4
## 2      Austria 199410     5
## 3      Austria 199512     5
## 4      Austria 199910     4
## 5      Austria 200211     5
## 6      Austria 200610     5
## 7      Austria 200809     6
## 8      Belgium 199111    11
## 9      Belgium 199505    10
## 10     Belgium 199906     9
```

Grouping operations: number of parties per election final

```r
# using "chaining" -- no need for intermediate objects
nparties2 <- distinct(cmpdataLong, countryname, date, party) %>%
    group_by(countryname, date) %>%
    summarise(npart = n())
identical(nparties, nparties2)

## [1] TRUE
```

Relational data bases

- ▸ invented by E. F. Codd at IBM in 1970
- ▸ A relational database is a collection of data organized as a set of formally defined tables
- ▸ These tables can be accessed or reassembled in many different ways without having to reorganize the underlying tables that organize the data
- ▸ RDBMS: a relational database management system. Examples include: MySQL, SQLite, PostgreSQL, Oracle. MS Access is a lite version of this too.
- ▸ The standard user and application programmer interface to a relational database is structured query language (SQL)

Example

- example: Database of Parties, Elections, and Governments (DPEG) relational database

```
SELECT c.countryName, c.countryAbbrev, p.* FROM party AS p
    LEFT JOIN country AS c
    ON p.countryID = c.countryID
```

- simpler example: convert CMP data into relational tables for countries, parties, elections, categories, and counts

Basic relational structures

- tables
  - also known as "relations"
  - tables define the forms of the data that are linked to other data through key relations
- keys: how tables are cross-referenced
  - primary key: an column in a table that uniquely identifies the remaining data in the table
  - foreign key: a field in a relational table that matches the primary key column of another table
  - `join` operations link tables in a structured query

Normal forms 1

"Normalizing" a database means creating a proper set of relations
First normal form: No Repeating Elements or Groups of Elements

```
head(select(cmpdata, countryname, partyname, date, per108, per110))

##      countryname                      partyname   date per108 per110
## 175     Austria             FP Freedom Party 199010      3      0
## 176     Austria          GA Green Alternative 199010      0      3
## 177     Austria SP Social Democratic Party 199010      5      0
## 178     Austria           VP People's Party 199010      8      0
## 179     Austria             FP Freedom Party 199410      1     11
## 180     Austria             LF Liberal Forum 199410      0      0
```

Here, this is violated because of the wide format of per108 and
per110. To solve this, we have to move this to long format.

Normal forms 2

Second normal form: No Partial Dependencies on a Concatenated Key

```
head(cmpdataLong)
```

```
##   countryname party   date category catcount
## 1     Austria 42420 199010   per101        0
## 2     Austria 42110 199010   per101        0
## 3     Austria 42320 199010   per101        0
## 4     Austria 42520 199010   per101        5
## 5     Austria 42420 199410   per101        0
## 6     Austria 42421 199410   per101        0
```

Here, the format is still violated, because party 42420 is repeated. To solve this we need to create a party table and link to it using a foreign key.

Normal forms 3

Third normal form: No Dependencies on Non-Key Attributes.
Every non-prime attribute of data in a table must be dependent on
a primary key.

```
head(cmpdataLong)

##   countryname party   date category catcount
## 1     Austria 42420 199010   per101        0
## 2     Austria 42110 199010   per101        0
## 3     Austria 42320 199010   per101        0
## 4     Austria 42520 199010   per101        5
## 5     Austria 42420 199410   per101        0
## 6     Austria 42421 199410   per101        0
```

Here, this is violated because election data repeats across multiple
values of the category count table, when it should have its own
table.

Non-relational data

- recently popularized because so much data is unstructured, and dealing with new data forms in a classic relational setting requires changing the entire schema
- non-relational systems typically define data using key-value pairs
- example: JSON - see
  `http://kenbenoit.net/files/JSONexample.json`

Compression: sparse matrix format

used because many forms of matrix are very sparse - for example, document-term matrixes

```
require(quanteda, warn.conflicts = FALSE, quietly = TRUE)
myDfm <- dfm(inaugTexts, verbose = FALSE)
myDfm[1:10, 1:5]

## Document-feature matrix of: 10 documents, 5 features.
## 10 x 5 sparse Matrix of class "dfmSparse"
##                  features
## docs              fellow-citizens  of the senate and
##    1789-Washington               1  71 116      1  48
##    1793-Washington               0  11  13      0   2
##    1797-Adams                    3 140 163      1 130
##    1801-Jefferson                2 104 130      0  81
##    1805-Jefferson                0 101 143      0  93
##    1809-Madison                  1  69 104      0  43
##    1813-Madison                  1  65 100      0  44
##    1817-Monroe                   5 164 275      0 122
##    1821-Monroe                   1 197 360      0 141
##    1825-Adams                    0 245 304      0 116
```

Compression: sparse matrix format

used because many forms of matrix are very sparse - for example,
document-term matrixes

```
# how many cell counts are zeros
sum(myDfm==0) / (ndoc(myDfm) * nfeature(myDfm)) * 100

## [1] 91.67571

object.size(myDfm)

## 1123544 bytes

object.size(as.matrix(myDfm))

## 4762040 bytes
```

Additional course preview

Day	Date	Topic(s)	Details
Mon	17 Aug	Course overview and introduction to data science	We will use this session to get to know the range of interests and experience students bring to the class, as well as to survey the approaches to be covered. We will also discuss and demonstrate the R software.
Tue	18 Aug	Research design issues in data science	Sampling, causal inference from observational data, differences with experimental settings, features. variables. Basic probability and statistics, binomial and Normal distributions. Cross-validation, predictive accuracy versus marginal effects. Model selection.
Wed	19 Aug	Linear Regression	The basic linear regression model, with a focus on prediction.
Thu	20 Aug	Generalized linear regression	Logistic regression, GAMs.
Fri	21 Aug	Resampling methods	Cross-validation, bootstrap.

Additional course preview

Mon	24 Aug	Association rules and clustering	Cluster analysis, k-means clustering, and hierarchical clustering
Tue	25 Aug	Machine Learning	Decision trees, k-Nearest Neighbour, Naive Bayes. Evaluation metrics, precision and recall, cross-validation.
Wed	26 Aug	Unsupervised learning and dimensional reduction.	Principal components analysis, correspondence analysis.
Thu	27 Aug	Text analysis	Working with text in R, sentiment analysis, dictionary methods.
Fri	28 Aug	Mining the Social Web	Working with the Twitter API, Facebook API, JSON data, and examples.

Preview of the first lab session

- ▶ Using the GitHub repository
- ▶ Editing and running the file in RStudio using RMarkdown
- ▶ Submitting the file to Moodle